

# Logistic regression

→ Classification problems

## Hyperparameters: (model performance)

- Penalty
- C
- L1\_ratio
- solver
- max\_iter
- tol
- warm\_start
- multiclass
- fit\_intercept
- intercept\_scaling
- class\_weight

## More hyperparameters:

(computational choices)

- Random\_state: int  
reproducibility
- Verbose: int  
Output level
- n-jobs: int  
parallel execution

TOTAL: 14 !

Penalty: str, default: l2

- l2: bias/variance tradeoff
- l1 :dimensionality reduction
- Elasticnet: combination of l1,l2
- None

C: float, default: 1.0

→ Grid search

l1\_ratio: float,  
→ Grid search      default: 1.0

Multiclass: str, default: auto

- ovr: a binary problem is fit for each label.
- multinomial: the loss minimised is the multinomial loss fit across the entire probability distribution, even when the data is binary.  
unavailable when solver=liblinear
- Auto: ovr if binary or liblinear, else multinomial

Solver: str,  
default: lbfgs

- liblinear
- lbfgs
- newton-cg
- sag
- saga

tol: float,  
default: 0.0001

max\_iter: int,  
default: 100

liblinear:

Library for Large Linear Classification.  
Moves toward the minimum in one direction at a time.  
→ “one-vs-rest” for multiclass

newton-cg:

Exact Hessian (2nd derivative)  
→ SLOW

lbfgs:

Limited-memory BFGS  
(Broyden–Fletcher–Goldfarb – Shanno).  
Approximates 2nd derivatives

sag:

Stochastic Average Gradient descent  
→ Fast, NEEDS SCALED DATA!!  
(reduce tol?)

saga:

Extension of sag that also allows for L1 regularization (the only one...).

[https://scikit-learn.org/dev/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/dev/modules/linear_model.html#logistic-regression)  
<https://towardsdatascience.com/dont-sweat-the-solver-stuff-aea7cddc3451>